

Group 14:

Parallelisation of graph algorithms in Julia

Speaking Order:

1. Blain Cribb
2. Aorthi Afroza
3. Kelsey Murray

Overview

Blain Cribb:

- Julia

Aorthi Afroza:

- Considered algorithms
- Prim's algorithm

Kelsey Murray:

- A* Algorithm
- Julia benchmarking tools



- High-level language
- General-purpose
 - Data science
 - Machine learning
 - Fast numerical and scientific computing
 - Parallel computing
- Dynamically typed but with great performance
- Open-sourced with fantastic documentation
- “Looks like Python, feels like lisp, runs like Fortran”

Julia - Looks like Python

- Code is easy to read and write
- Designed to be easy to pick up for other language users
- Plenty of modern features that make coding easier:
 - Large base library
 - Garbage collection
 - Optional data typing
 - Interoperability

Julia - Feels like Lisp

- Support for metaprogramming
 - Code is stored as a data structure within the language
 - This code can be manipulated during runtime
 - Very useful for applications such as machine learning
- True Lisp-style macros unlike static macros
- Multiple dispatch for OOP like behaviour

Julia - Runs like Fortran

- Created to be fast from its inception
- Static analysis and just in time compilation allows for running of some code in a static manner
- Optional static typing means with smart programming we can achieve static-like speeds
- Multiple dispatch can help with the static analysis

Julia - Runs like Fortran



<https://julialang.org/benchmarks/>

Julia - Parallelism

- Natively supported
- Multithreading (experimental)
- Loops are easily parallelisable through use of macros
- Supports atomic access for variables

Considered Algorithms


- Criteria for algorithms:
 - Different algorithm types; e.g. shortest path, traversal, minimum spanning tree
 - Algorithms with a considerable amount of literature out there
 - Successful existing implementations of their parallelised algorithms

Julia - LightGraphs.jl


- A Julia Package that provides the framework for:
 - Building graphs
 - Traversing them
 - Building your own graphing algorithms
- Contains the module Parallel, which contains parallel implementations of:
 - Bellman Ford Shortest Paths
 - Dijkstra Shortest Paths
 - Floyd Warshall Shortest Paths
 - Johnson Shortest Paths
 - Bfs
 - Greedy Color
 - 5 Centrality measures
 - 4 Distance measures

Prim's Algorithm

- Minimum spanning tree (MST)
 - A tree that connects all nodes in the graph
 - The **least total cost** among all trees that connect all the nodes
 - This may not be the shortest path

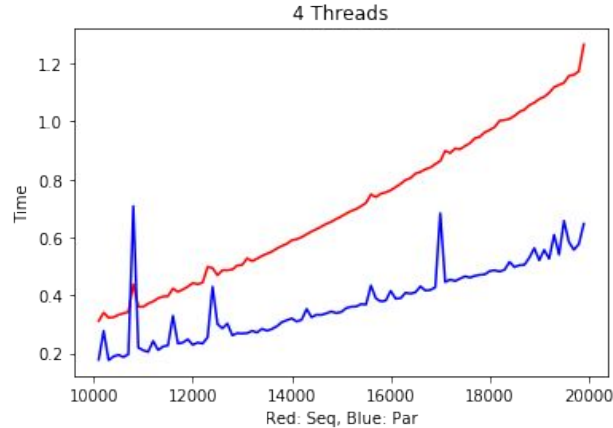


Prim's Algorithm: Example



Previous Attempts at **Parallelising** Prim's


- PES University in India
- University of California, Irvine
 - Implemented in OpenMP




<https://github.com/parthvshah/parallel-prim>

- University of Singapore
 - Implemented in C++
 - This paper proposes an adaptation of Prim's Algorithm

Implementation Plan: Prim's



	a	b	c	d	e	f
d[i]	1	0	5	1	∞	∞
a	0	1	3	∞	∞	2
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	5	5	0



	a	b	c	d	e	f
d[i]	1	0	2	1	4	∞
a	0	1	3	∞	∞	2
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	5	5	0





Figure 7.6 The partitioning of the distance array d and the adjacency matrix A among p processors. Copyright (r) 1994 Benjamin/Cummings Publishing Co.

Testing and Benchmarking: **Prim's**


- Datasets:
 - Densely connected, large, undirected graphs
- Metrics
 - Speed-up
 - Correctness




A* Algorithm

- **Pathfinding algorithm**
 - Popular in gaming where there's one source, one destination, and possibly obstacles
 - “Informed” Dijkstra's algorithm
 - Cost of a considered path takes into account heuristic as well as edge weight
- **Best-first search, but not greedy**
 - Doesn't solely rely on what seems best at that time (the heuristic) but also includes the exact cost of the potential path (the edge weights)

A* Algorithm: Example



Dijkstra's Algorithm





A* Algorithm

Parallel A* in Literature

First Approach:



- Looks at entire current best **path**
- Investigates all the nodes surrounding the **path**



Parallel A* in Literature


Second Approach:

- Looks at the current **end node** of the current best path
- Investigates nodes directly adjacent to that **end node**




Implementation Plan: A*


- Can we move **diagonally**? Will this make a difference to the speedup?
 - If we can move diagonally: **Euclidean** distance as heuristic
 - If we can't: **Manhattan** distance as heuristic



Manhattan





Euclidean



Testing and Benchmarking: A*

- Metrics:
 - Speedup
 - Correctness: Do we care if the **exact path** is the same?
 - Ideally parallel is **at least as good** as serial but we don't care if the path isn't exactly the same
- Datasets:
 - Large grid maps with obstacles



Julia **Benchmarking** Tools

- @time macro

Output:

```
0.244729 seconds (294.16 k allocations: 14.614 MiB, 5.74% gc time)
```

- @benchmark macro

Output:

```
BenchmarkTools.Trial:
  memory estimate: 2.13 KiB
  allocs estimate: 19
  -----
  minimum time:      1.770 μs (0.00% GC)
  median time:       2.170 μs (0.00% GC)
  mean time:         3.924 μs (37.21% GC)
  maximum time:      9.772 ms (99.92% GC)
  -----
  samples:           10000
  evals/sample:      10
```

Thank you!

